

Attacker Cognitive Loop: An Interpretable Rule-Based Engine for Estimating, Predicting, and Engaging Adaptive Attackers

Jiahao Lai
ASSCOR Project
chins-xing@proton.me

August 20, 2026

Abstract

Attackers and defenders operate under fundamentally unequal information endowments: attackers choose when, where, and how to probe; defenders react to signals that arrive too late and too sparse. This paper presents the Attacker Cognitive Loop (ACL), a system designed to actively reverse this information asymmetry — not by blocking attacker actions, but by guiding them through lightweight decoys that reveal intent earlier than the attacker intends. ACL closes the loop between observation and intervention, turning the defender’s information disadvantage into a controllable orchestration problem. The engine is implemented as an interpretable, rule-based extension of the open-source ASSCOR platform, with every component mapped to a public repository path and branch. Its heuristic weights are explicitly not optimal; they are provided solely to demonstrate closed-loop mechanics and to serve as a baseline for community-driven replacement. We formulate the research question — *how AI, shared tooling, and individual experience jointly shape the attacker’s action distribution, and whether defenders can actively reshape it through observation, deception, and guidance* — and state ten open research questions (Q1–Q10) with evidence ratings. We provide the complete formal specification of the scoring, normalization, sharpness, and utility functions, a reproducible four-round closed-loop trace with exact numerical values, a sensitivity analysis of the hand-tuned weights, and simulated per-round latency benchmarks.

1 Introduction

Classical defense embeds blocking controls: access control and isolation make undesirable states unreachable. Surveys of moving-target defense (MTD) have documented two inherent weaknesses of this paradigm. First, a *cognitive limitation*: defenders cannot enumerate all latent attacker paths, and attackers enjoy an information advantage during prolonged reconnaissance. Second, a *temporal asymmetry*: attackers can plan over long horizons while defenders must respond in real time, and patch lag leaves a window of opportunity.

At the heart of the two weaknesses above lies a single structural fact: attackers hold an inherent information advantage. They observe the defender’s environment at their chosen pace, select which paths to probe, and can wait for the defender’s response patterns to emerge. Defenders, by contrast, see only the artifacts of actions already taken — alerts, logs, and blocked connections. This is not a technical failure of firewalls or IPS; it is a cognitive asymmetry built into the defender’s observational posture.

Guided active defense, as we define it, is an attempt to reverse this asymmetry. Instead of waiting for the attacker to reveal intent through damage, the defender proactively deploys ephemeral decoys that appear to reduce the attacker’s uncertainty — inducing the attacker to verify hypotheses early, thereby exposing TTPs and intent before the attack reaches its objective. ACL is designed as an orchestration engine for exploring this reversal: it estimates what the attacker believes, predicts what the attacker will do next, and selects decoys that maximize information gain while minimizing the defender’s exposure. The current implementation provides a fully auditable rule-based baseline against which more principled models can be compared.

Guided active defense does not aim to block everything; it aims to create an *information asymmetry* in the defender’s favor. By deploying lightweight decoy targets and fabricated intelligence, the defender induces the attacker to abandon slow, patient operation in favor of rapid verification — exposing tactics, techniques, and procedures (TTPs) and intent earlier than they would otherwise be revealed. Intelligence collected from each interaction feeds the next, more accurate round of deception, forming a positive feedback loop: *guide* $\rightarrow$ *collect* $\rightarrow$ *analyze* $\rightarrow$ *decide* $\rightarrow$ *guide again*.

In the MITRE Engage taxonomy, engagement is used here as a guidance tool, not a blocking tool; host isolation is demoted to a fallback when guidance fails or confirms high risk.

This work makes the following contributions:

1. We formulate the attacker-cognitive-loop research question that connects attacker modeling with adaptive deception.
2. We present an interpretable, rule-based engine (ACL) that realizes the loop: state estimation, action-distribution prediction, utility-driven engagement selection, and sharpness-driven strategy switching, with the complete formal specification of every score, normalization, sharpness, and utility term (§4.3.1, §4.5.1, §4.4.1).
3. We provide a reproducible four-round closed-loop trace with exact numerical values computed from the reference implementation’s weights (§9), so the engine’s behavior can be independently verified and replayed.
4. We anchor the work in a real, public codebase: the engine is an implemented extension of the open-source ASSCOR assessment platform, with every component mapped to a repository path and branch (§7, §4), so the paper is directly auditable rather than a paper-only design.
5. We state ten open research questions (Q1–Q10) about attacker behavior in the AI era, each with an evidence rating, and a validation plan.
6. We are transparent about what is implemented (the engine) versus what remains to be answered (the questions), and we analyze the sensitivity of the hand-tuned weights (§12.1). No assumption is presented as a conclusion.

2 Related Work

Decision-theoretic planning. Shinde and Doshi’s I-POMDP² [1] extends partially observable Markov decision processes to model the defender’s beliefs about the attacker’s beliefs, capabilities, and preferences, with finite nested recursive reasoning. Their experiments observe level-3 defenders who first mislead the attacker into believing the defense is passive, then deploy decoys — a decision-theoretic justification for guidance over blocking. A full I-POMDP² solver is computationally prohibitive; our engine is a transparent rule-based approximation that can be replaced by richer models without changing the loop architecture.

Game-theoretic deception. Zhu’s tutorial [2] and Anwar & Kamhoua’s attack-graph games [5] formalize signaling games, dynamic games, mechanism design, and hypergames for deception. General-sum stochastic games make pure-strategy Nash equilibrium computation PSPACE-hard, and equilibrium concepts assume symmetric optimality that poorly models the attacker’s *bounded rationality*. We therefore prefer the decision-theoretic view.

Attacker modeling and attribution. The Diamond Model of intrusion analysis [6] structures every intrusion event as four vertices — *adversary*, *capability*, *infrastructure*, and *victim* — connected by six phases of an activity thread, with meta-features (timestamp, kill-chain phase, result, direction, method, resources) describing each event. Its central axiom, that each intrusion is an event in an activity thread whose evolution is governed by the adversary, is the analytical backbone we reuse: our six-layer graph (§4.1) is the Diamond Model’s relational projection onto a temporal substrate — the Adversary layer holds the adversary vertex, the Capability layer holds capabilities and TTPs, the Network/Dependency layers hold infrastructure, and the Identity/Evidence layers hold the victim and the meta-features. Unlike the Diamond Model, which is a retrospective analytic framework, ACL consumes the same relational structure *forward* to estimate the attacker’s state and predict its next action. We also note the fundamental attribution problem raised by the shared-capability questions (Q5/Q6, §10): observed similarity between intrusions does not imply the same adversary, since adversaries share tooling, code, and (increasingly) AI-generated strategy.

Defensive deception and honeypot limitations. Zhang and Thing [3] survey three decades of honeypots, honeytokens, and MTD across a four-layer stack (network/system/software/data) mapped to the kill chain, and document the two structural weaknesses of honeypot-only defenses that motivate our guided-engagement design: (i) *static deployment*: if honeypots and honeytokens are left with static configuration, an adversary has enough time to infer their existence, map them, and evade them; and (ii) *pivot risk*: high-interaction honeypots, which offer a real operating-system environment, can themselves be compromised and used as a pivot point to attack other systems. Two further weaknesses are well-known in practice: (iii) *low-interaction honeypots have low fidelity*, so sophisticated attackers can fingerprint and bypass them; and (iv) *passivity*: traditional honeypots wait for the attacker and yield little intelligence about intent until the attacker has already committed. Our loop addresses (i) by making engagement a *function of the predicted action distribution* that changes each round, (ii) by the *sufficiency principle* (§4.4.2) — decoys expose only a plausible break-in appearance with no real system to compromise — (iii) by keeping decoys lightweight and cheap enough to deploy at single-host scale, and (iv) by using decoys as *active sensors* whose interaction chains feed the state engine.

MTD design principles. MTD surveys [4] contribute five design principles (coverage, unpredictability, timeliness, superstability, functional equivalence) that constrain any future MTD component. We adopt honeypot/honeytoken as the primary guidance mechanism and treat MTD as an optional deterministic-weakening component.

Behavior standardization. MITRE ATT&CK [7] provides the semantic layer for attack behavior. Our prediction engine maps action distributions onto TTP distributions rather than replacing ATT&CK, and our state engine infers intent from observed TTPs through an explicit technique→intent table. Table 1 lists the mapping used by the reference implementation (the full table is §4.2); the same table drives the reverse projection $P(TTP_i) = \sum_j P(TTP_i | Action_j) P(Action_j)$ (§4.3.1). This makes the ATT&CK→state-machine bridge explicit, auditable, and replaceable: unknown techniques remain **unknown** rather than being guessed.

Risk propagation. SRD (Systemic Risk Dynamics) models risk diffusion over a network graph [9]; we reuse its substrate for the topology layer and as the dynamic state-evolution layer of the cognitive loop.

ATT&CK technique	Inferred intent	Example
T1595, T1592, T1590, T1046	recon	Active Scanning (T1595)
T1110, T1003, T1555, T1078	credential	Brute Force (T1110)
T1021, T1570, T1550, T1028	lateral	Remote Services (T1021)
T1567, T1048, T1030, T1537	data_theft	Exfiltration Over Web Service (T1567)
T1190, T1193, T1566	web_attack	Exploit Public-Facing Application (T1190)

Table 1: TTP $\rightarrow$ intent mapping used by the state engine (§4.2). Techniques outside this table remain **unknown**; the mapping is open data, not engine logic.

Our work is built on ASSCOR (ASSess + CORE), an open-source security acceptability assessment platform. ASSCOR provides the three substrates that ACL requires: (i) a six-layer temporal graph (Network, Dependency, Identity, Attacker, Capability, Evidence) that carries observations and alerts as temporal events; (ii) SSAM 2.0, a multidimensional computable assessment model that produces per-host acceptability scores $\sigma \in [0, 100]$; and (iii) Prism/SRD, a three-layer risk propagation engine that models how compromise diffuses through the network. Without these substrates, ACL’s cognitive loop would have no data to consume — the state engine needs graph topology to localize observations, the predictor needs SSAM scores as vulnerability signals, and the engagement planner needs Prism risk edges to prioritize containment targets.

ACL is thus not a replacement for these engines; it is an orchestration layer that closes the loop between assessment (passive) and intervention (active). ASSCOR’s main branch provides the stable assessment baseline; ACL lives in the ASSCOR-Research-Core branch, an experimental sandbox explicitly designated for active-defense prototyping. This split is not merely a code-management convenience — it reflects the different development rhythms of assessment (stable, auditable) and active defense (iterative, replaceable), and ensures that experimental prediction logic never compromises the production-grade assessment pipeline.

A central design boundary further separates ACL from learning-based approaches. Neural-network-based predictors have demonstrated impressive performance in low-stakes security tasks such as log correlation, anomaly scoring, and alert triage. However, ACL operates in a regime where individual recommendations carry operational risk: a decoy deployed on the wrong target, or at the wrong time, can alert the attacker, waste defensive resources, or be pivoted against. In this regime, interpretability under failure is a safety requirement, not a desirable extra. When a defensive intervention fails, the operator must be able to trace the decision to specific, auditable factors — not to aggregate model weights. ACL therefore defaults to a fully decomposable rule-based scorer (Eq. 6). The architecture remains open to neural replacements at Seam 2 (§6.3), with the caveat that any proposed replacement must preserve a mechanism for post-hoc explanation (e.g., through surrogate models or attention-based attribution) so that the safety requirement is not violated.

3 Problem Formulation

3.1 Attacker State

We model the attacker’s cognitive state as:

$$\text{AttackerState} = (\text{Capability}, \text{Experience}, \text{Intent}, \text{TTP_Repertoire}, \text{SharedCapability}, \\ \text{IndividualCapability}, \text{AI_Dependence}, \text{TargetKnowledge}, \text{BeliefState}, \text{Objective})$$

where $Intent \in \{\text{RECON}, \text{CREDENTIAL}, \text{LATERAL}, \text{DATA_THEFT}, \text{WEB_ATTACK}\}$, $AI_Dependence \in [0, 1]$ measures reliance on AI for decision/execution, and $BeliefState$ captures the attacker’s subjective perception of the target environment.

3.2 Prediction Target

Rather than a deterministic next-TTP, we predict:

$$P(A_{t+1} \mid AttackerState_t, TargetState_t, Observation_t)$$

and additionally output *most likely*, *most dangerous*, and *most observable* actions (and a ranked list), feeding both risk prioritization and engagement selection.

3.3 The Cognitive Loop

The closed loop is:

$$\begin{aligned} Observation &\rightarrow State\ Estimation \rightarrow Behavior\ Prediction \rightarrow Guidance/Deception \\ &\rightarrow Attacker\ Response \rightarrow New\ Evidence \rightarrow State\ Update \rightarrow Predict\ Again \end{aligned}$$

4 System Design: The Attacker Cognitive Loop

ACL is implemented as a modular, pure-function library (no external dependencies per component), designed so each component can be independently tested and replaced. It follows a microkernel architecture: the core kernel remains untouched; these components are optional, build-tag modules.

4.1 Multi-Layer Temporal Graph

A single `Host--connects--Host` graph cannot express attackers, identities, tools, code, TTPs, and evidence. We maintain six semantic layers — Network, Dependency, Identity, Attacker, Capability, Evidence — over a shared temporal graph: nodes and edges carry *FirstSeen/LastSeen* timestamps, edges carry lifecycle status (up/down) and attributes (latency, bandwidth, ACL), and nodes carry an event stream (observations, alerts, decoy triggers) with confidence. Reachability queries (shortest path, bounded-hop traversal) support multi-hop analysis. The layer/type/edge-type vocabularies are open strings: predefined constants cover the common semantics, and any model can introduce new values without engine changes.

4.2 Attacker State Engine

The state engine implements $OldState + Evidence \rightarrow NewState$ with explicit, interpretable rules. Formally, let the state at round t be $s_t = (K_t, E_t, I_t, \mathcal{R}_t, \mathcal{C}_t, B_t)$ where K_t is target knowledge, E_t experience, I_t intent, $\mathcal{R}_t$ the TTP repertoire, $\mathcal{C}_t$ the capability set, and B_t the belief map. An evidence item $e = (ttp, intent, target, outcome, c)$ with confidence $c \in [0, 1]$ updates the state by:

$$K_{t+1} = \min(1, K_t + 0.1 \cdot c \cdot \mathbf{1}[target \neq \emptyset]) \quad (1)$$

$$E_{t+1} = \min(1, \max(0, E_t + 0.05 \mathbf{1}[outcome = \text{success}] - 0.03 \mathbf{1}[outcome = \text{failure}])) \quad (2)$$

$$B_{t+1}[target] = \min(1, B_t[target] + 0.15 \cdot c) \quad (3)$$

$$\mathcal{R}_{t+1} = \mathcal{R}_t \cup \{ttp\}, \quad \mathcal{C}_{t+1} = \mathcal{C}_t \cup \{cap(ttp)\} \quad (4)$$

$$I_{t+1} = \arg \max_{e'} c_{e'} \quad \text{over evidence with known intent,} \quad (5)$$

where explicit *intent* overrides $IntentFromTTP(ttp)$

where $\text{cap}(\cdot)$ maps a technique to its capability requirement. Updates are pure: the input state is never mutated. Intent inference is deliberately conservative — unknown TTPs remain **unknown** rather than being guessed (see Algorithm 1 in §8.1).

4.3 Prediction Engine

4.3.1 Raw Scoring

The prediction engine converts state plus target state into a normalized action distribution over six actions $a \in \mathcal{A} = \{\text{recon}, \text{credential}, \text{lateral}, \text{data_theft}, \text{web_attack}, \text{maintain}\}$. Let the state at round t be s_t with intent I_t , experience E_t , target knowledge K_t , and AI dependence A_t ; and let the target have SSAM score $\sigma \in [0, 100]$ and exposure $\rho \in [0, 1]$. The raw score of action a is the sum of six interpretable terms:

$$\begin{aligned} \text{score}_t(a) = & b_a + w_{\text{int}} \mathbf{1}[I(a) = I_t] + E_t \mathbf{1}[a \in \mathcal{A}_{\text{exec}}] - 0.5 E_t \mathbf{1}[a = \text{recon}] \\ & + K_t \mathbf{1}[a \in \mathcal{A}_{\text{kn}}] + A_t \mathbf{1}[a \in \mathcal{A}_{\text{ai}}] - 0.5 A_t \mathbf{1}[a = \text{recon}] \\ & + w_{\text{vuln}} v_t \mathbf{1}[a \in \mathcal{A}_{\text{vuln}}]. \end{aligned} \quad (6)$$

with

$$v_t = \frac{100 - \sigma}{100} + \rho \quad (7)$$

and the action subsets are

$$\begin{aligned} \mathcal{A}_{\text{exec}} &= \{\text{lateral}, \text{data_theft}\}, & \mathcal{A}_{\text{kn}} &= \{\text{credential}, \text{lateral}, \text{data_theft}\}, \\ \mathcal{A}_{\text{ai}} &= \{\text{credential}, \text{web_attack}\}, & \mathcal{A}_{\text{vuln}} &= \{\text{credential}, \text{lateral}, \text{data_theft}, \text{web_attack}\}. \end{aligned}$$

The baseline priors and weights used by the reference implementation are listed in Table 2. All scores are clamped to $[0, \infty)$ before normalization.

Parameter	Default value
b_a (baseline prior)	recon 1.0, credential 1.2, lateral 1.1, data_theft 0.9, web_attack 1.0, maintain 0.5
w_{int} (intent continuity)	3.0
w_{vuln} (vulnerability weight)	0.8
T (softmax temperature)	1.0
danger weight (MostDangerous)	data_theft 1.0, lateral 0.8, credential 0.6, web_attack 0.5, recon 0.3, maintain 0.2
$\alpha, \beta, \gamma, \delta$ (utility)	1.0, 0.8, 0.6, 1.2
sharpness threshold θ	0.3

Table 2: Default weights and priors of the reference implementation (Table 6, 13). All values are configuration, not engine logic.

4.3.2 Normalization

Scores are converted to a probability distribution by a temperature- controlled softmax (numerically stabilized by subtracting the maximum score):

$$P_t(a) = \frac{\exp((\text{score}_t(a) - \max_{a'} \text{score}_t(a'))/T)}{\sum_{a'' \in \mathcal{A}} \exp((\text{score}_t(a'') - \max_{a'''} \text{score}_t(a'''))/T)} \quad (8)$$

which satisfies $P_t(a) \geq 0$ and $\sum_{a \in \mathcal{A}} P_t(a) = 1$. The distribution is then projected onto the ATT&CK TTP layer:

$$P(TTP_i) = \sum_{a \in \mathcal{A}} P(TTP_i | a) P_t(a), \quad P(TTP_i | a) = 1/|\mathcal{T}_a| \text{ for } TTP_i \in \mathcal{T}_a, \quad (9)$$

where $\mathcal{T}_a$ is the technique set associated with action a (Table 1 and its reverse). This keeps ATT&CK as the behavioral semantic layer rather than replacing it.

4.3.3 Multi-Outputs

Beyond the full distribution, the engine outputs:

$$a_t^{\text{ML}} = \arg \max_a P_t(a) \quad (10)$$

$$a_t^{\text{MD}} = \arg \max_a P_t(a) d_a \quad (\text{danger-weighted}) \quad (11)$$

$$a_t^{\text{MO}} = \arg \max_{a \in \{\text{recon}, \text{credential}, \text{lateral}\}} P_t(a) \quad (\text{most observable}) \quad (12)$$

with danger weights d_a from Table 2, feeding risk prioritization and engagement selection.

4.4 Engagement Planner

4.4.1 Utility Function

The engagement planner selects deceptive interventions by an explicit utility function. For an intervention E targeting action a_E with deployment cost $cost_E$ and exposure exp_E , the utility is:

$$\text{Utility}(E) = \alpha \underbrace{P_t(a_E) \kappa(\sigma, \rho)}_{\text{IG}(E)} + \beta \underbrace{DP_E}_{\text{detection probability}} + \gamma \underbrace{AV_E}_{\text{attribution value}} - \delta \underbrace{cost_E \cdot exp_E}_{\text{Risk}(E)} \quad (13)$$

with the target-coverage factor

$$\kappa(\sigma, \rho) = 0.5 + 0.5 \cdot \frac{(100 - \sigma)/100 + \rho}{2}. \quad (14)$$

Information gain $\text{IG}(E)$ is the predicted probability of the action the decoy targets, scaled by coverage; DP_E and AV_E are intrinsic decoy capabilities (Table 3); Risk is the product of cost and exposure. Only interventions with positive utility are candidates; they are returned in descending utility order.

Decoy	Target action	DP	AV	Cost/Exposure
fake SSH	lateral	0.90	0.60	0.30 / 0.20
fake credential	credential	0.80	0.90	0.20 / 0.15
fake document	data_theft	0.75	0.85	0.20 / 0.15
fake web	web_attack	0.70	0.50	0.40 / 0.30
scan port	recon	0.85	0.30	0.10 / 0.10

Table 3: Default decoy catalog: detection probability (DP), attribution value (AV), and cost/exposure per decoy type.

4.4.2 Decoy as Sensor

Following the *sufficiency principle*, decoys need only create a plausible break-in appearance, not high-fidelity honeypot credentials. Decoys are treated as sensors: an interaction chain (connect $\rightarrow$ credential attempt $\rightarrow$ command $\rightarrow$ system discovery $\rightarrow$ follow-up) is recorded as a *DeceptionRecord* and converted into an *Evidence* item (with the TTP/intent inferred from the interaction chain, conservatively: only when a TTP is identified) that feeds the state engine, closing the loop (§3.3).

4.4.3 Decoy Anti-Pivot and Self-Destruct

Because the engagement planner operates as a sensor, its decoys must not become a second attack surface: an adversary who compromises a decoy could otherwise pivot from it onto production assets. All deployed decoys are therefore **ephemeral** (TTL ≤ 60 s). If a decoy receives any *outbound* connection request, it immediately terminates the session and resets to a clean state. The decoy environment is **air-gapped from production assets**: it holds no real credentials, no routing tables, and no reachable references to the production network. Combined with the sufficiency principle (§4.4.2), this guarantees that a decoy can collect at most one bounded interaction chain before it disappears, and that nothing reachable from a decoy can be used to reach production. The reference implementation enforces this contract in the decoy deployer (lifecycle timer + outbound-connection teardown); the corresponding defensive-fallback strategy (§4.5.2) further ensures that containment, not decoy deployment, is the default response to unknown intent or unknown TTPs.

4.5 Closed-Loop Controller

4.5.1 Decision Sharpness

We use normalized entropy as a sharpness metric to trigger strategy switching; it is not calibrated accuracy. Let

$$S_t = 1 - \frac{H(P_t)}{\ln |\mathcal{A}|}, \quad H(P_t) = - \sum_{a \in \mathcal{A}} P_t(a) \ln P_t(a), \quad (15)$$

so sharp distributions (low entropy) have high sharpness and flat ones low. **This metric is a heuristic for action-distribution peakedness, not a probabilistic guarantee of correctness.** A flat distribution — the honest expression of an ambiguous attacker — yields low sharpness, which drives the controller toward containment or collection rather than aggressive engagement; the strategy switch (16) never interprets low entropy as high certainty.

4.5.2 Strategy

The controller selects a strategy by first applying a defensive fallback and then thresholding sharpness:

$$\text{strategy}_t = \begin{cases} \text{contain} & I_t = \text{UNKNOWN or } \exists \text{ unknown TTP in observations} \\ \text{collect} & S_t < \theta \quad (\text{keep only the highest-IG intervention: intelligence first}) \\ \text{engage} & S_t \geq \theta \quad (\text{keep all positive-utility interventions: active guidance}) \end{cases} \quad (16)$$

with default $\theta = 0.3$. **Contain** denotes temporary network isolation or process suspension of the affected target until new evidence arrives; no decoys are deployed while containing. It guarantees that the worst case — an attacker using a technique absent from the TTP $\rightarrow$ intent mapping, or an intent that cannot be established — leaves the system in a safe state instead of a passively

open one. This realizes the information-gain closed loop: predict $\rightarrow$ select $\rightarrow$ observe $\rightarrow$ update $\rightarrow$ re-predict.

5 Design Rationale

This section documents the concrete design choices and the trade-offs they entail, so that the system can be evaluated and modified by others.

5.1 Rule-Based over Learned or Game-Theoretic Models

The state engine and predictor use hand-crafted rules and weights. This choice reflects the operational constraints of security decision-making, not a temporary expedient.

First, accountability under failure. In security, the cost of a wrong decision is asymmetric and operationally consequential. A decoy deployed against the wrong attacker intent does not merely reduce accuracy; it may signal the defender’s posture to the attacker, or waste the limited window in which the attacker is observable. When such a failure occurs, the operator needs to know why. A neural predictor, by itself, cannot answer that question without post-hoc explanation tools, which introduce their own approximation errors and are not yet standardized for audit trails. ACL’s rule-based scorer, by contrast, emits a decomposition of every score term at prediction time (Eq. 6), so every decision can be inspected, audited, and, if necessary, corrected. This is a non-negotiable requirement for a system that recommends actions with real operational impact.

Second, no training data. The research questions Q1–Q10 are not yet answered; there is no labeled dataset of attacker behavior from which to learn. Rules are the only honest starting point. But even with abundant data, we would not replace the scoring function with a black-box model without preserving full decomposability — the safety requirement is structural, not data-dependent.

Third, computational cost. Full I-POMDP² or general-sum stochastic-game solution is computationally prohibitive in real-time constraints. A transparent approximation is acceptable as a first implementation, and the loop architecture is deliberately model-agnostic so that better white-box approximations can replace the current rules without sacrificing auditability.

We emphasize: this is a constitutional principle of the ASSCOR platform, not a preference. ACL is designed to prioritize determinism and traceability over raw predictive accuracy. The three replacement seams (§6.3) are explicitly provided so that the community can substitute more principled models — but any replacement must preserve explainability, or it violates the safety contract of the system.

5.2 Intent-Continuity Weight

The action whose intent equals the current *AttackerState.Intent* receives a strong additive weight ($w_{\text{int}} = 3.0$, versus baseline 0.5–1.2). Rationale: the current stage is the strongest single predictor of the next action — an attacker currently harvesting credentials tends to continue credential actions. The magnitude was tuned so that, under a clear intent, the distribution becomes sharp enough for the sharpness-driven controller to switch to active engagement, while remaining responsive to new contradictory evidence. The sensitivity of this choice is analyzed in §12.1.

5.3 Entropy-Based Sharpness

Decision sharpness is defined as the normalized entropy of the action distribution (Eq. 15): sharp distributions have high sharpness, flat ones low. It is a *peakedness heuristic*, not calibrated accuracy

— we state this explicitly in §12. The default threshold ($\theta = 0.3$) separates the collect strategy (flat distributions: prefer intelligence gathering) from the engage strategy (sharp distributions: allow active guidance), and the defensive fallback (§4.5.2) overrides both whenever intent is unknown or an unknown TTP is observed. Both the threshold and the temperature are configuration parameters.

5.4 Utility Weights

The engagement utility (Eq. 13) uses default weights $\alpha = 1.0$, $\beta = 0.8$, $\gamma = 0.6$, $\delta = 1.2$. Information gain dominates because intelligence acquisition is the primary objective of the loop (§3.3); risk carries the largest coefficient because decoy exposure (detection as fake, or collateral business impact) is the main cost of deception. These weights are tunable per deployment.

5.5 Open Semantic Vocabulary

Layer, *NodeType*, and *EdgeType* are plain strings. Predefined constants cover the common semantics (six graph layers, 23 node types, nine edge types), but any model can introduce new values without changing the engine. This decouples the graph engine from any specific domain vocabulary — the same engine serves network topology, identity, attacker, capability, and evidence layers.

5.6 Pure-Function, Immutable Updates

Update (state) and *Predict* (distribution) never mutate their inputs; they return new values. This yields three properties: components are trivially unit-testable; concurrent evaluation is safe; and the full state/prediction trajectory can be retained for audit and replay.

5.7 Decoy Mapping Reuse

The default catalog reuses the existing intent→decoy mapping of the platform (Table 3): lateral movement → fake SSH, credential theft → fake credentials, data theft → fake documents, web attack → fake web, scanning → decoy ports. This follows the *sufficiency principle*: decoys need only present a plausible break-in appearance, not high-fidelity honeypot credentials, which keeps deployment cost low enough for single-host-scale operation.

6 Model Integration

6.1 Microkernel, Optional Modules

ACL follows the host platform’s microkernel architecture: the kernel retains only the extension framework and interface contracts; every ACL component is an optional, independently compiled module (build-tag gated). The kernel is not modified — the loop consumes data through existing interfaces (topology registry, assessment results) and publishes evidence back.

6.2 Data-Flow Wiring

Three data paths feed the loop:

- **Evidence in:** observations, alerts, CTI, and decoy interaction chains are normalized to *Evidence*{*Source*, *TTP*, *Intent*, *Target*, *Outcome*, *Confidence*} and fed to the state engine.

- **Target state:** the platform’s assessment pipeline provides target properties (SSAM score, exposure, zone) via *TargetState*.
- **Evidence out:** decoy interaction chains are recorded as *DeceptionRecord* (connect → credential attempt → command → system discovery → follow-up) and converted to *Evidence*, closing the loop.

6.3 Model Replacement Points

The loop is model-agnostic at three narrow seams. Each seam is a single interface with one method; the graph layer and the closed-loop controller depend only on these interfaces, never on the internal choices of any engine.

Seam 1 — StateEngine. The rule-based belief update (Eq. 1–5) can be replaced by any model exposing *Update(state, evidence) → state*. Replacement examples: (a) a Bayesian belief-update model (e.g., Beta-Bernoulli per technique) that maintains posteriors over intent conditioned on evidence, then collapses them into the same *AttackerState* fields; (b) a wrapper around an I-POMDP² solver that projects its belief tree onto the *AttackerState* structure after each update.

Seam 2 — Predictor. The rule-based scorer (Eq. 6–8) can be replaced by any calibrated probabilistic model exposing *Predict(state, target) → distribution*. Replacement examples: (a) a logistic-regression or gradient-boosted classifier over engineered state features, calibrated (e.g., Platt scaling) so the output is a genuine probability distribution over the six actions; (b) a neural sequence model conditioned on the evidence history, with the two-layer ATT&CK projection (Eq. 9) retained downstream.

Seam 3 — Planner. The utility function (Eq. 13) and decoy catalog are configuration; richer planners can be swapped in behind the same *Select(distribution, target) → interventions* interface. Replacement examples: (a) an information-theoretic optimal experimental design that selects the decoy minimizing expected posterior entropy of the attacker state; (b) a reinforcement-learning planner whose reward is evidence gain minus deployment cost.

In every case the controller calls the same three methods (*Update* → *Predict* → *Select*) and the loop contract (§3.3) is preserved. The reference implementation ships the rule-based engines as the default implementations of these interfaces.

6.4 Open Vocabulary Integration

Because the graph semantics are open strings (§5.5), a new model (e.g., an identity layer feeding an ID-theft detection model) can register its own layers and edge types and immediately participate in temporal queries and reachability analysis without forking the engine.

7 Reference Implementation and Repository Context

This section anchors the paper in its actual codebase. Every formula and algorithm above is implemented, unit-tested, and runnable today in a public repository; no component of the loop exists only on paper.

7.1 Repository and Branches

The reference implementation lives in the public repository <https://github.com/chins-xing/Asscor> (module `github.com/asscor/asscor`, Go 1.26, Apache-2.0). Two long-lived branches exist:

- **main** — the stable baseline (v0.2.x): the security acceptability assessment runtime (SSAM/PRISM/SRD engines, microkernel, plugin system) without the attack-loop extension.
- **ASSCOR-Research-Core** — the topology-specialized research branch (formerly v0.3.0) on which this paper’s ACL engine and its adversarial modules are developed and maintained. The paper’s claimed implementation artifacts are reproducible from this branch at the referenced commit.

The branch split is deliberate: **main** preserves the stable, production-facing assessment platform, while the research branch carries experimental topology and engagement work without destabilizing the release baseline.

7.2 Host Platform

ACL does not stand alone: it is an optional, build-tag-gated extension of the ASSCOR security acceptability assessment runtime, an open-source evaluation platform (Apache-2.0) that continuously assesses the security acceptability of each monitored host. The platform provides the loop’s inputs and consumes its outputs through existing interfaces:

- **Assessment pipeline:** SSAM 2.0 (system security acceptability model) computes per-host acceptability scores $\sigma \in [0, 100]$; the Prism/SRD engine (Systemic Risk Dynamics) propagates risk over the network graph. Both engines are published as standalone zero-dependency Go modules (github.com/chins-xing/ssam, github.com/chins-xing/prism) and vendored in `ssam-lib/` and `prism-lib/`.
- **Microkernel and plugins:** the kernel keeps only the extension framework (dependency injection, event bus, extension registry, plugin lifecycle); all 17 functional modules are optional build-tag plugins. ACL components are modules in the same regime.
- **Reproducible lab:** a containerized multi-host lab (Containerlab, 18 nodes) is the network substrate for Phase 1 validation (§11).

7.3 ACL Component Mapping

Table 4 maps every ACL component to its implementation location in the repository, so a reader can locate, build, and audit the exact code behind each formula.

Component	Location	Roles
Attacker state engine	<code>internal/attackerstate/</code>	state update Eq. 1–5, TTP→intent/capability tables
Action predictor	<code>internal/predictor/</code>	scoring Eq. 6, softmax Eq. 8, TTP projection
Engagement planner	<code>internal/engagement/</code>	utility Eq. 13, decoy catalog, decoy→evidence sensor
Closed-loop controller	<code>internal/defensecycle/</code>	Step, sharpness Eq. 15, three-way strategy Eq. 16 incl. contain
Intent-driven deception	<code>optional/adversary/packages/mitre-engage/</code>	lightweight decoys (ports/tokens), anti-pivot lifecycle
Reproduction harness	<code>cmd/tracecheck/</code> (build tag <code>tracecheck</code>)	replays §9 and §12.1 numbers

Table 4: ACL component → repository path mapping. All listed packages compile and pass their unit tests in the **ASSCOR-Research-Core** branch.

7.4 Reproducibility

Three layers guarantee that the paper’s numbers are checkable:

1. **Unit tests:** the four engine packages ship with tests covering intent inference, distribution normalization, utility decomposition, containment behavior, and multi-round convergence.
2. **Trace reproduction:** `cmd/tracecheck` (invoked as `go run -tags tracecheck` with `./cmd/tracecheck`) recomputes the four-round trace (§9) and both sensitivity tables (§12.1) directly from the engine code, so every value in this paper can be regenerated in seconds.
3. **Lab environment:** the Containerlab multi-host topology (`docs/clab-lab/`) plus the SSAM/PRISM assessment pipeline constitute the reproducible substrate for the Phase 1 validation experiments.

No number in this paper is claimed from a private or unverifiable run; the repository commit referenced in this paper is the single source of truth.

8 Key Implementations

The engine is implemented as pure-function libraries. The following condensed pseudocode captures the core of each component; the formal specifications are in §4.2–4.5.2.

8.1 State Update

```
Algorithm 1: Update(state, evidence list) -> newState
for each ev in evidence:
  if ev.TTP is known: add ev.TTP to state.TTPRepertoire
  if ev.Target != "": state.TargetKnowledge += 0.1 * clamp01(ev.Confidence)
  switch ev.Outcome:
    case "success": state.Experience += 0.05
    case "failure": state.Experience -= 0.03
  intent := ev.Intent if set, else IntentFromTTP(ev.TTP)
  if intent known and ev.Confidence >= bestConf: // high-conf wins
    state.Intent, bestConf = intent, ev.Confidence
  if capability = CapabilityForTTP(ev.TTP): add to state.Capability
  if ev.Target != "": state.BeliefState[ev.Target] += 0.15 * conf
clamp Experience, TargetKnowledge, BeliefState to [0,1]
return state
```

Key property: no mutation of the input — the new state is a fresh value.

8.2 Action Distribution

```
Algorithm 2: Predict(state, target) -> distribution
for each action a in {recon, credential, lateral, data_theft, web, maintain}:
  score[a] = base[a]
    + 3.0 if intentOf(a) == state.Intent // continuity
    + (a in {lateral,data_theft}) * state.Experience
    - (a == recon) * 0.5 * state.Experience
```

```

    + (a in {credential,lateral,data_theft}) * state.TargetKnowledge
    + (a in {credential,web}) * state.AIDependence
    - (a == recon) * 0.5 * state.AIDependence
    + (a in {credential,lateral,data_theft,web})
      * vulnerability * 0.8
distribution = softmax(score, temperature)    // Eq. (13)
mostDangerous = argmax_a distribution[a] * dangerWeight[a]
return distribution with multi-outputs

```

8.3 Engagement Utility

Algorithm 3: Select(distribution, target) -> scored interventions

```

for each intervention E (decoy type, target action):
    IG = distribution[E.targetAction] * coverage(target)    // Eq. (9)
    DP = detectionRate[E.decoy]    // intrinsic decoy capability
    AV = attributionValue[E.decoy]
    Risk = E.cost * E.exposure
    utility = alpha*IG + beta*DP + gamma*AV - delta*Risk    // Eq. (8)
return interventions sorted by utility (descending), utility > 0

```

8.4 Closed-Loop Step

Algorithm 4: Step(state, observations, target) -> result

```

newState      = StateEngine.Update(state, observations)
distribution   = Predictor.Predict(newState, target)
sharpness     = 1 - entropy(distribution) / log(|actions|)    // Eq. (15)
strategy      = "contain" if intent unknown or unknown TTP observed
               = "collect" if sharpness < threshold
               = "engage" otherwise    // Eq. (16)
candidates    = Planner.Select(distribution, target)
interventions = strategy == "contain" ? []    // isolate
               : strategy == "collect" ? [argmax_IG(candidates)]    // intel
               : candidates    // guide
return {newState, distribution, sharpness, strategy, interventions}

```

All four algorithms are deterministic and side-effect free; the reference implementation ships with unit tests covering intent inference, distribution normalization, utility decomposition, and multi-round convergence to a credential intent with the strategy switching to engage (§9).

9 Illustrative Closed-Loop Trace

To make the engine’s behavior concrete and independently verifiable, this section replays a four-round closed-loop run against a single target, computed exactly from the default parameters in Table 2 and the reference implementation’s formulas (Eq. 1–16). The scenario follows the engine’s multi-round convergence test: the attacker is observed scanning the target, then engages in credential attacks against it; the defender’s decoys trigger and feed new evidence back into the loop.

Scenario. Target **host-a** with SSAM score $\sigma = 55$ and exposure $\rho = 0.6$; vulnerability factor $v = (100 - 55)/100 + 0.6 = 1.05$; coverage factor $\kappa = 0.5 + 0.5 \cdot 1.05/2 = 0.7625$. Initial attacker state: unknown intent, zero experience, zero target knowledge.

Round 1 (no observation yet). The state engine has no evidence, so $I = \text{unknown}$, $E = 0$, $K = 0$, $A = 0$. A subtle but important engine behavior appears here: because the continuity check matches $I(a) = I_t$ and the reference implementation maps **maintain** to intent **unknown** (“no intent yet, keep status quo”), the intent-continuity term (ii) in Eq. 6 applies to **maintain** as well, giving it +3.0. Scores from Eq. 6:

action a	recon	credential	lateral	data_theft	web_attack	maintain
b_a	1.0	1.2	1.1	0.9	1.0	0.5
vulnerability term ($0.8v$)	–	+0.84	+0.84	+0.84	+0.84	–
continuity term (w_{int})	–	–	–	–	–	+3.0
score(a)	1.00	2.04	1.94	1.74	1.84	3.50

Softmax (Eq. 8, $T=1$), subtracting the max score 3.50:

action	recon	credential	lateral	data_theft	web_attack	maintain
$\exp((s - \max)/T)$	0.082	0.232	0.210	0.172	0.190	1.000
$P(a)$	0.044	0.123	0.111	0.091	0.101	0.530

Entropy $H = 1.425$; sharpness $S = 1 - 1.425/\ln 6 = 0.205$. Although $S < \theta$, the defensive fallback (§4.5.2) takes precedence: the intent is still **unknown**, so the strategy = **contain** — no decoys are deployed and the target is kept isolated until the first observation arrives.

Round 2 (attacker scans). Observation: T1595 (Active Scanning), target **host-a**, confidence 0.6. State update (Eq. 1–5): $K = 0.06$, $B[\text{host-a}] = 0.09$, $I = \text{recon}$ (T1595 $\rightarrow$ recon), repertoire {T1595}. Scores:

action	recon	credential	lateral	data_theft	web_attack	maintain
score(a)	4.00	2.10	2.00	1.80	1.84	0.50
$P(a)$	0.649	0.097	0.088	0.072	0.075	0.020

(recon receives +3.0 intent-continuity and +0.06 knowledge; execution actions receive +0.06 knowledge and +0.84 vulnerability.) Entropy $H = 1.181$; $S = 1 - 1.181/\ln 6 = 0.341 \geq \theta$; intent is now known, so the fallback does not trigger and strategy = **engage**. All five decoys have positive utility; the ranking is **scan port** ($U = 1.343$), **fake credential** ($U = 1.218$), **fake document** ($U = 1.129$), **fake SSH** ($U = 1.075$), **fake web** ($U = 0.773$), so all are deployed.

Round 3 (decoy triggered: brute force). The fake credential decoy records a brute-force interaction: T1110, outcome success, confidence 0.85. State update: $E = 0.05$, $K = 0.145$, $B = 0.2175$; $I = \text{credential}$ (T1110 $\rightarrow$ credential, confidence 0.85 > previous 0.6); repertoire {T1595, T1110}. Scores:

action	recon	credential	lateral	data_theft	web_attack	maintain
score(a)	0.975	5.185	2.135	1.935	1.84	0.50
$P(a)$	0.013	0.873	0.041	0.034	0.031	0.008

Entropy $H = 0.567$; $S = 1 - 0.567/\ln 6 = 0.683$; strategy = **engage**. The distribution is now sharply peaked at **credential**; $IG(\text{fake credential}) = 0.873 \times 0.7625 = 0.666$, the top-ranked intervention.

Round 4 (credential dumping). The decoy records T1003 (OS Credential Dumping), success, confidence 0.9. State update: $E = 0.10$, $K = 0.235$, I remains **credential**. Scores:

action	recon	credential	lateral	data_theft	web_attack	maintain
score(a)	0.95	5.275	2.275	2.075	1.84	0.50
$P(a)$	0.012	0.874	0.044	0.036	0.028	0.007

Entropy $H = 0.561$; $S = 0.687$; strategy = **engage**. The intent has converged to **credential**, sharpness has risen from 0.205 to 0.687, and the strategy has moved from defensive containment (unknown intent) through intelligence-first collection to active guidance (engage) once the distribution became sharp enough.

Table 5 summarizes the trajectory. All values above can be reproduced exactly by executing the reference implementation with the default parameters (Table 2); the engine’s unit tests assert the same convergence behavior (intent $\rightarrow$ credential, strategy $\rightarrow$ engage).

Round	Observation / trigger	Intent	E	K	S	Strategy
1	(none)	unknown	0.00	0.00	0.205	contain
2	T1595 scan (conf 0.6)	recon	0.00	0.06	0.341	engage
3	T1110 brute force, success (conf 0.85)	credential	0.05	0.145	0.683	engage
4	T1003 dumping, success (conf 0.9)	credential	0.10	0.235	0.687	engage

Table 5: Four-round closed-loop trajectory: intent converges to **credential**, sharpness rises $0.205 \rightarrow 0.687$, and the strategy moves from contain (unknown intent) through engage as evidence accumulates. Exact distributions per round are in the text.

9.1 Performance Benchmark (Simulated)

To substantiate the real-time requirement (§1), we report simulated single-round latency of the reference implementation as a function of topology scale. The engine’s per-round work is dominated by three steps: state update (§8.1), prediction (§8.2), and engagement selection (§8.3); graph-layer reachability queries scale with the topology size. Table 6 shows simulated wall-clock timings on a single commodity host (the simulation injects synthetic graph sizes and measures the three engine steps; values are indicative, not production benchmarks).

Topology size (nodes)	State update (ms)	Prediction (ms)	Engagement (ms)	Total loop (ms)
50	0.8	1.2	0.5	2.5
500	12.4	15.1	6.2	33.7
5000	185.0	210.3	89.5	484.8

Table 6: Simulated per-round latency vs. topology size (single host, indicative values).

For sub-second response requirements, the engine is suitable for networks up to 500 nodes without optimization; larger deployments require distributed sharding (left as future work).

10 Open Research Questions and Working Assumptions

We state ten open research questions about attacker behavior, together with the working assumptions our engine currently embodies. Evidence ratings follow a five-star scale (★★★★ = direct empirical support, ★ = speculative). We explicitly distinguish established theory from inference requiring validation. **This paper does not claim to answer these questions; it provides the measurement framework required to answer them.**

ID	Question	Status	Evidence
Q1	Does successful experience raise the reuse probability of related TTPs?	To be validated	★★★☆☆
Q2a	Does IT skill positively correlate with hacking efficiency?	Empirically supported	★★★★★
Q2b	Do capability/experience raise reliance on validated TTPs?	Inference	★★★☆☆
Q3	Do attackers seek initial advantage (identity/trust/information gain)?	Multi-source support	★★★★☆
Q4	Can social-engineering strategic value be measured by frequency alone?	Methodological	★★★★☆
Q5	Does a shared capability layer exist (common tools/code/AI)?	Direct evidence	★★★★★
Q6	Does AI expand the accessibility of shared capability?	Strong support	★★★★★
Q7	Does AI lead to attacker strategy convergence?	Core unvalidated question	★★☆☆☆
Q8	Are effective capability and strategy diversity independent?	Theoretical inference	★★☆☆☆
Q9	Can high experience offset AI convergence?	To be validated	★★☆☆☆
Q10	Should one predict action distributions rather than next TTP?	Methodological support	★★★★☆

Table 7: Open research questions and evidence ratings. The engine’s rules embody tentative answers (working assumptions) only where evidence exists; all other questions remain open.

Implications. Q3 implies attackers first reduce target uncertainty rather than directly exploiting vulnerabilities — guiding through decoys that appear to reduce that uncertainty is therefore aligned with attacker incentives. Q5/Q6 imply that observed similarity does not imply the same actor: attribution must first exclude shared tooling, shared code, common TTPs, shared infrastructure, and AI-generated strategy — this is exactly the Diamond Model’s attribution caveat (§2). Q7 is the core research question — if AI dependence raises cross-attacker similarity, defenders can exploit the resulting predictability.

We explicitly do *not* pre-commit to conclusions such as “AI always lowers attack entropy” or “a specific probability model is always better”; these are to be tested.

11 Validation Plan

Validation is planned in two phases.

Phase 1 — engine validation (in progress). The engine is implemented and unit-tested (intent inference, distribution normalization, utility decomposition, loop convergence in simulation; §9 shows an exact replay). A containerized multi-host lab (Containerlab, 18 nodes) serves as the network substrate; the SSAM/PRISM assessment pipeline provides target-state inputs (scores, topology, propagation edges). The next step is a scripted multi-host attack scenario (injection of realistic TTP sequences) whose closed-loop traces are recorded and published as part of the reproducible artifact.

Phase 2 — question validation (planned). Each question maps to a measurable experiment:

- **Q1/Q9:** longitudinal study of a red-team exercise — measure TTP reuse conditioned on prior success/failure and operator experience.

- **Q2a/Q2b:** controlled human-subject study correlating IT skill with efficiency and TTP reliance.
- **Q7:** comparative analysis of AI-assisted vs. human-only attackers — TTP entropy, sequence diversity, behavioral distance, strategy cluster distribution (this is the core, currently undated, question).
- **Q10:** ablation comparing next-TTP prediction vs. action-distribution prediction on the same dataset.

We note that the 2026 DBIR [8] provides empirical grounding for the threat landscape (exploitation as top initial-access vector at 31%, credential abuse, ransomware prevalence 48%, GenAI-assisted attacks with a median of 15 techniques), but dataset sampling/reporting biases require preserving source and provenance metadata rather than treating frequency as ground truth.

12 Discussion

Limitations. ACL currently uses hand-crafted rules and weights; it is a transparent approximation, not a full I-POMDP² solver. Decision sharpness is defined as distribution peakedness (Eq. 15), not calibrated accuracy — it is a heuristic for triggering strategy switches, not a probabilistic guarantee of correctness (§4.5.1). Deception utility weights $(\alpha, \beta, \gamma, \delta)$ require tuning. The research questions Q1, Q2b, Q7, Q9 are not yet answered; the paper does not claim predictive accuracy or deception success without direct evidence. The four-round trace (§9) demonstrates *mechanical* convergence of the engine in simulation; it is *not* evidence about real attackers.

12.1 Weight Sensitivity

The engine’s behavior is governed by hand-tuned weights, so we analyze how sensitive the outputs are to the two most consequential parameters: the intent-continuity weight w_{int} and the softmax temperature T .

Intent-continuity weight. Holding all other parameters at their defaults (credential intent, $E = 0.05$, $K = 0.145$, target as in §9), Table 8 shows $P(\text{credential})$ and the sharpness S as w_{int} varies. The action distribution is highly sensitive: raising w_{int} from 0 to 4.5 moves $P(\text{credential})$ from 0.255 to 0.969, and sharpness from 0.07 to 0.90. The default 3.0 sits in the steepest part of the curve, meaning small re-tunings materially change how quickly the controller switches to engage: with $w_{\text{int}} = 1.5$ the sharpness just crosses the threshold ($S = 0.302$), while with $w_{\text{int}} = 0$ the distribution stays flat enough that the loop remains in collect (intent is known, so the fallback does not apply).

w_{int}	0.0	1.5	3.0 (default)	4.5
$P(\text{credential})$	0.255	0.605	0.873	0.969
sharpness S	0.07	0.30	0.68	0.90
strategy	collect	engage	engage	engage

Table 8: Sensitivity of the predicted distribution and sharpness to the intent-continuity weight w_{int} (credential intent, target as in §9); values verified against the reference implementation.

Temperature. The softmax temperature T controls distribution sharpness directly. Table 9 varies T at round 2 of the trace (recon intent, scores $\{4.00, 2.10, 2.00, 1.80, 1.84, 0.50\}$). At $T = 0.5$

the distribution is near-deterministic ($P(\text{recon}) = 0.937$, $S = 0.82$); at $T = 2$ it flattens enough that sharpness falls below the threshold and the strategy reverts to collect even though the attacker’s intent is already established. Temperature and the sharpness threshold θ must therefore be treated as a joint calibration pair, not independently.

T	0.5	1.0 (default)	2.0
$P(\text{recon})$	0.937	0.649	0.384
H (entropy)	0.324	1.181	1.637
sharpness S	0.82	0.34	0.086
strategy	engage	engage	collect

Table 9: Sensitivity of distribution sharpness and strategy to softmax temperature T (round 2 of the trace in §9, recon intent, scores $\{4.00, 2.10, 2.00, 1.80, 1.84, 0.50\}$).

Implication for scientific value. These sensitivities are exactly why the rule-based predictors are positioned as *first approximations* rather than calibrated models (§5): the current weights encode plausible heuristics, but the outputs are scale-sensitive, and any deployment must either calibrate the weights against real attacker data (Phase 2, §11) or replace the scorer with a calibrated model at Seam 2 (§6.3). The formal specification (§4.3.1–4.5.2) exists precisely so that such calibration and replacement can be done and evaluated without reverse-engineering the implementation.

Relationship to richer models. The loop architecture is model-agnostic: the rule-based state engine and predictor can be replaced by learned models or an I-POMDP² solver without changing the graph, engagement, or controller layers. This modularity is a deliberate design property.

The AI convergence question. If Q7 is answered affirmatively, defenders face both a threat and an opportunity: convergent strategies are easier to predict and to guide; diverse strategies (Q8/Q9) resist convergence and require broader engagement portfolios.

13 Conclusion

We presented the Attacker Cognitive Loop, an interpretable rule-based engine for guided active defense that estimates attacker state, predicts action distributions, and selects utility-driven deceptive interventions within a closed loop that switches between containment, information gathering, and active engagement by decision sharpness. We gave the complete formal specification of the scoring, normalization, sharpness, and utility functions (§4.3.1–4.5.2), a reproducible four-round closed-loop trace with exact values (§9), a sensitivity analysis of the hand-tuned weights (§12.1), and simulated per-round latency benchmarks (§9.1). The engine is not a paper-only design: it is fully implemented as a modular, unit-tested extension of the open-source ASSCOR platform, with every component located in the public repository (§7, §4). We stated ten open research questions with evidence ratings and a validation plan, and we were explicit about the boundary between implementation and unvalidated assumption. The primary contribution of this work is the interface contract and reproducible simulation harness. We do not claim defensive efficacy; instead, we offer the community a transparent, modular testbed for falsifying or validating the stated research questions.

The fundamental problem ACL addresses is not the inadequacy of any particular defensive control; it is the structural information asymmetry between attacker and defender. By closing the loop between observation, prediction, and intervention, ACL provides a systematic mechanism to

reverse that asymmetry — not by making the defender omniscient, but by making the attacker’s uncertainty costly to maintain. We release ACL as a modular, reproducible baseline, and invite the community to replace its heuristic components with more principled models. The ultimate measure of success is not the accuracy of any single prediction, but the defender’s ability to act before the attacker’s information advantage becomes decisive.

Because the three engine seams (§6.3) are narrow interfaces, the community can evaluate better models against this open-source, reproducible baseline. Contact: `chins-xing@proton.me`.

Ethics Statement

This work is exclusively intended for **legitimate defensive purposes**. All capabilities described — attacker state estimation, action prediction, and deceptive engagement — are designed to be deployed only by authorized defenders against adversarial activity within systems and networks they are legally permitted to protect. The authors do not provide or endorse any offensive capability, and we assume all research, experimentation, and validation is conducted in authorized environments in compliance with applicable laws, regulations, and institutional policies. The reference implementation is distributed solely as a defensive security research artifact.

References

- [1] A. Shinde and P. Doshi. “Decision-theoretic planning and cognitive modeling for active cyber deception.” *Artificial Intelligence*, vol. 356, article 104540, 2026. <https://doi.org/10.1016/j.artint.2026.104540>
- [2] Q. Zhu. “Game Theory for Cyber Deception: A Tutorial.” In *Proceedings of the 2019 Hot Topics in the Science of Security Symposium and Bootcamp (HoTSoS ’19)*, Nashville, TN, USA, April 2019, ACM. arXiv:1903.01442. <https://arxiv.org/abs/1903.01442>
- [3] L. Zhang and V. L. L. Thing. “Three Decades of Deception Techniques in Active Cyber Defense — Retrospect and Outlook.” *Computers & Security*, vol. 106, article 102288, 2021. <https://doi.org/10.1016/j.cose.2021.102288> Also arXiv:2104.03594.
- [4] C. Lei, H.-Q. Zhang, J.-L. Tan, Y.-C. Zhang, and X.-H. Liu. “Moving Target Defense Techniques: A Survey.” *Security and Communication Networks*, vol. 2018, article ID 3759626, 2018. <https://doi.org/10.1155/2018/3759626>
- [5] A. H. Anwar and C. Kamhoua. “Game Theory on Attack Graph for Cyber Deception.” In *Decision and Game Theory for Security (GameSec 2020)*, LNCS, Springer, 2020. https://doi.org/10.1007/978-3-030-64793-3_24
- [6] S. Caltagirone, A. Pendergast, and C. Betz. “The Diamond Model of Intrusion Analysis.” Technical report, Center for Cyber Intelligence Analysis and Threat Research (CCIATR), 2013.
- [7] MITRE. “MITRE ATT&CK.” <https://attack.mitre.org> (accessed 2026)
- [8] Verizon. *2026 Data Breach Investigations Report*. 2026. <https://www.verizon.com/business/resources/reports/dbir/>
- [9] ASSCOR Project. “Active Defense Design Whitepaper” (v0.3.0 draft), 2026. Internal document: attacker model (H1–H10), automated attack/defense chain, deception taxonomy and MTD survey summaries; reference implementation of the ACL engine in `internal/attackerstate`, `internal/predictor`, `internal/engagement`, `internal/defensecycle`. The trace and sensitivity tables of this paper (§9, §12.1) are reproduced by `cmd/tracecheck` (build tag `tracecheck`) against the reference implementation. Repository: <https://github.com/chins-xing/Asscor> (branch ASSCOR-Research-Core, Go 1.26, Apache-2.0).